

Final Report for STAT 451/551

Jack Barney, Kyle Maher, Max Donelan, & Tate Wietfeld

Spring 2026

Introduction

The given FP data contains 200 observations of eight predictors labeled x_1 through x_8 , and the resulting response parameter y . Additional consideration was given to predicting y_{bin} which is 1 if y is greater than or equal to 70 and 0 otherwise. x_1 through x_7 are continuous while x_8 is categorical with four groups A, B, C , and D . No context surrounding the predictors or response parameters was given. This report is segmented into initial Exploratory Data Analysis, followed by the Regression Models, then the Classification Models, and concludes with some overall observations. All written code can be found in the appendix.

Exploratory Data Analysis

Continuous y

Correlations




Figure 1: Correlations of Continuous Parameters

Pairplot




Figure 2: Scatterplots of Continuous Parameters

y




Figure 3: Histogram of y




Figure 4: Histograms of y by x_8

- A right skewed
- D left skewed

Dichotomized y

Regression Models

Plotting the 10-Fold Cross Validation Metrics




Figure 1: RMSE by Model for Each Fold




Figure 2: RMSE by Model and by Fold




Figure 3: MAE by Model for Each Fold

Classification Models

Standard Binary Logistic Regression

To fit the standard logistic regression, the first step we need to take is to redefine the dependent variable. Previously, y was able to take on any integer value from 0 to 100, but in the case defined by this project, y would need to be reduced to a binary variable depending on its previous larger scale values. This new version of y , y_{bin} , is given a value of 1 if y 's previous value was greater than or equal to 70, or 0 otherwise. With variable y_{bin} defined, our next step is to fit the first version of our standard logistic regression model.

Using the `glm` function in R, we can fit a standard logistic regression using every predictor variable at our disposal (x_1 through x_8) in efforts to model y_{bin} . The following Figure 1 provides the numerical summary of that model:

```
Call:
glm(formula = ybin ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8, family = binomial,
    data = FP_di)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -2.3626      0.5001  -4.724 2.31e-06 ***
x1             0.5494      0.2742   2.004 0.045122 *
x2            -0.1866      0.1575  -1.185 0.235959
x3            -0.1923      0.7719  -0.249 0.803300
x4             1.7534      1.6982   1.032 0.301851
x5            -0.7956      1.4050  -0.566 0.571203
x6             0.2261      0.5576   0.406 0.685090
x7             0.9472      0.2656   3.566 0.000363 ***
x8B           -1.7534      0.7377  -2.377 0.017461 *
x8C           -0.5180      0.7104  -0.729 0.465833
x8D           -1.8050      0.8173  -2.208 0.027211 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Figure 1: Summary of Initial Logistic Model

Immediately the lack of significant predictors is apparent, with the majority of the independent variables holding p-values greater than 0.05. At this point, we need to check whether we have violated the assumptions of a standard logistic regression. Those assumptions are as follows: The outcome variable is binary, the data points in each group are independent (little to no clustering), the independent variables must be linearly related to the log-odds of the outcome, there must be little to no multicollinearity, no drastic outliers, and finally that we have a large sample size (Leung, 2025). A few of these we can check off quickly: Our outcome variable `ybin` is binary, it's defined as described earlier in this section. Next, we can confirm that we have a large sample size because our data has 200 observations, which is generally considered large.

With that, we can begin to critique the other requirements. Starting with checking for outliers and clustering, we can plot the individual predictor variables against the original `y` to check for both. We use the original `y` because Those plots are as follows:




Figure 2: Binned y values vs individual predictors

Seen in Figure 2, none of the values from the predictor variables have extravagant outliers, nor do they cluster dramatically in any scatterplot. At the bottom right of Figure 2 we have the plot of ybin vs x8, which is different from the rest because x8 is a categorical variable. There's no outliers to be seen here either, because values of x8 can only be A, B, C, or D. Thus, we've met two more of our assumptions.

The next assumption we need to check is that our independent variables need to be linearly related

to the log odds of the outcome. This can be done using the Box-Tidwell test. To do so, we must first define a new variable for each predictor variable (x_n) that is equal to $x_n * \log(x_n)$. Then, we define a logistic regression model, modeling y_{in} with x_n and $x_n * \log(x_n)$ as the only predictors. When we look at the model's summary, if the $x_n * \log(x_n)$'s p-value is significant (< 0.05), then we cannot assume x_n is linearly related to the log odds of the outcome (Leung, 2025). These steps are repeated for each predictor variable, as shown in the below figure. Note that $newx_n$ represents the $x_n * \log(x_n)$ value for the n th predictor variable.

Call:

glm(formula = ybin ~ x1 + newx1, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.34223	0.86140	-1.558	0.119
x1	0.04347	1.05418	0.041	0.967
newx1	0.54778	1.02683	0.533	0.594

Call:

glm(formula = ybin ~ x2 + newx2, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.0057	1.0001	-1.006	0.315
x2	-0.9035	1.1473	-0.788	0.431
newx2	0.6501	0.8763	0.742	0.458

Call:

glm(formula = ybin ~ x3 + newx3, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.737	1.032	-2.654	0.00796 **
x3	2.083	1.206	1.727	0.08412 .
newx3	-1.400	1.115	-1.256	0.20911

Call:

glm(formula = ybin ~ x4 + newx4, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.9964	0.9628	-2.074	0.0381 *
x4	1.1803	1.1458	1.030	0.3030
newx4	0.6285	1.1704	0.537	0.5913

Call:

glm(formula = ybin ~ x5 + newx5, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.489	1.477	-1.008	0.313
x5	-3.005	1.869	-1.608	0.108
newx5	-1.347	4.089	-0.329	0.742

Call:

glm(formula = ybin ~ x6 + newx6, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.9941	0.9691	-3.090	0.0020 **
x6	1.2975	1.0345	1.254	0.2098
newx6	-5.1561	2.3349	-2.208	0.0272 *

Call:

glm(formula = ybin ~ x7 + newx7, family = binomial, data = FP_di)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.5459	0.8301	-0.658	0.51077
x7	-2.0577	1.0659	-1.931	0.05354 .
newx7	3.4799	1.1221	3.101	0.00193 **

Figure 3: Summaries of Box-Tidwell Models

Above in Figure 3, we notice that for predictors x_6 and x_7 , we have to reject our assumption that they are linearly related to the log-odds of y_{bin} due to their low p-values. This means that for the actual model, we must transform them slightly instead of simply using them as they come from the dataset. In doing so, we are able to meet our assumption of linearity with the log odds of the outcome.

To resolve the final assumption about standard logistic regression models, that there is little to no multicollinearity between predictors, we first need to create a model using all of our predictors. The

only exception to this model is x6 and x7, which will be transformed slightly due to our findings earlier. Upon doing so, we are left with the following model:

```
Call:
glm(formula = ybin ~ x1 + x2 + x3 + x4 + x5 + log(x6 + 3) + log(x7 + 3) + x8, family = binomial, data = FP_di)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -5.5560      1.6262  -3.417 0.000634 ***
x1              0.5522      0.2683   2.058 0.039555 *
x2             -0.2131      0.1545  -1.379 0.167953
x3             -0.1409      0.7558  -0.186 0.852110
x4              1.7530      1.6602   1.056 0.290992
x5             -0.7173      1.3766  -0.521 0.602309
log(x6 + 3)    0.9907      1.5041   0.659 0.510111
log(x7 + 3)    2.2197      0.6878   3.227 0.001250 **
x8B            -1.5674      0.7053  -2.222 0.026252 *
x8C            -0.4724      0.7020  -0.673 0.501025
x8D            -1.5639      0.7714  -2.027 0.042637 *
```

Figure 4: Improved Standard Logistic Regression Model Summarized

The serious lack of significant predictors is worrying, but a potential solution to that comes from dealing with multicollinearity within the model, if it exists. We can check using the “Cars” package in R, followed by the vif() function, with the above model as the input to it. Doing so garners this result:

	GVIF	Df	GVIF^(1/(2*Df))
x1	1.136026	1	1.065845
x2	1.109377	1	1.053270
x3	11.479255	1	3.388105
x4	27.360025	1	5.230681
x5	17.578727	1	4.192699
log(x6 + 3)	1.378928	1	1.174278
log(x7 + 3)	1.911865	1	1.382702
x8	1.533878	3	1.073903

Figure 5: Initial VIF Report

Due to the categorical nature of x8, we don’t have exact VIF values, and instead must use the adjusted VIF scores, denoted in Figure 5 as column headed by $GVIF^{1/(2 \cdot Df)}$. Adjusted VIF uses the same scale as the normal VIF but is compatible with categorical as well as numeric variables. As seen in Figure 5, there aren’t any adjusted VIF measurements above 10, but the lack of significant variables noted in Figure 4 implies that there still may be some issues.

If we drop the highest adjusted VIF variable from the model, x4, and recreate the model, we get the following report from function `vif()`:

	GVIF	Df	GVIF ^{1/(2*Df)}
x1	1.117639	1	1.057184
x2	1.094885	1	1.046367
x3	1.257233	1	1.121264
x5	1.470021	1	1.212444
log(x6 + 3)	1.379700	1	1.174606
log(x7 + 3)	1.928150	1	1.388579
x8	1.504567	3	1.070455

Figure 6: Improved VIF Report

This is a significant improvement from the previous model, implying minute amounts of multicollinearity. As a result of the nature of multicollinearity, any models with fewer predictors than the model represented in Figure 6 will also have the same or lower levels of multicollinearity. Therefore, our model has kept all the standard assumptions of standard logistic regression intact, and our final standard logistic regression model is as seen in Figure 7:

```
Call:
glm(formula = ybin ~ x1 + x2 + x3 + x5 + log(x6 + 3) + log(x7 +
3) + x8, family = binomial, data = FP_di)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  -5.5091     1.6268  -3.386 0.000708 ***
x1             0.5339     0.2623   2.036 0.041757 *
x2            -0.2057     0.1530  -1.344 0.178817
x3             0.6214     0.2489   2.496 0.012551 *
x5            -2.1350     0.3967  -5.382 7.37e-08 ***
log(x6 + 3)    0.8931     1.5025   0.594 0.552246
log(x7 + 3)    2.2867     0.6810   3.358 0.000786 ***
x8B           -1.6467     0.7050  -2.336 0.019500 *
x8C           -0.4669     0.6988  -0.668 0.504031
x8D           -1.5198     0.7598  -2.000 0.045481 *
```

Figure 7: Summary of Final Standard Logistic Regression Model

This model results in an 87.5% prediction accuracy, leaving room for improvement while still being quite decent at framing the data correctly.

Regularized Logistic Regression

This version of a logistic regression model adds in a penalty term that works to diminish overfitting. In order to find the error term, we needed to turn the dataset into a matrix and use the `cv.glmnet()` function to use cross-validation and get a model for potential values. For our error term we used the minimum lambda described by the cross-validation model and implemented it using Lasso (L1) regularization via the `glmnet()` function using our data in matrix form again. The resulting model returned an accuracy of 88%, which is barely better than that of the standard logistic regression model.

Logistic GAM

Known also as the logistic generalized additive model, is a method for categorizing your outcome variable when you suspect that its predictors don't follow a straight line (Daniel, 2025). We designate such variables by putting them inside the parenthesis of `s()`, and using the modeling function `gam()` from the "mgcv" package. Considering its similar setup to standard logistic regression, we chose to use the same predictor variables as we did in that model, but using the `s()` on every predictor except for `x8`, because `x8` is categorical. This led to a model seen in with an accuracy of 95.5%, a fantastic improvement upon the original logistic regression model.

<pre>ybin ~ s(x1) + s(x2) + s(x3) + s(x5) + s(log(x6 + 3)) + s(log(x7 + 3)) + x8</pre>					Approximate significance of smooth terms:				
Parametric coefficients:						edf	Ref.df	Chi.sq	p-value
	Estimate	Std. Error	z value	Pr(> z)	s(x1)	2.669	3.316	6.999	0.10236
(Intercept)	-4.149	1.914	-2.167	0.03021 *	s(x2)	1.000	1.000	3.589	0.05817 .
x8B	-1.398	1.151	-1.214	0.22479	s(x3)	8.834	8.978	15.206	0.08415 .
x8C	-0.514	1.173	-0.438	0.66126	s(x5)	4.869	5.709	19.226	0.00344 **
x8D	-3.373	1.302	-2.592	0.00955 **	s(log(x6 + 3))	1.000	1.000	0.001	0.97595
---					s(log(x7 + 3))	6.065	7.161	19.449	0.00777 **
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1									

Figure 8: Logistic GAM Model Summary

Plotting the 10-Fold Cross Validation Metrics




Figure 9: Accuracy by Model for Each Fold




Figure 10: False Positive Rate by Model for Each Fold




Figure 11: False Negative Rate by Model for Each Fold




Figure 12: Accuracy by Model and by Fold

References

Leung, K. (2025, March 5). Assumptions of logistic regression, clearly explained. Towards Data Science. <https://towardsdatascience.com/assumptions-of-logistic-regression-clearly-explained-44d85a22b290/>

Daniel. (2025). Understanding Generalized Additive Models. RPubs. <https://rpubs.com/AfroLogicInsect/GAM>